

# 01\_traffic\_volume\_eda\_and\_baseline\_regression

July 27, 2025

## 1 Urban Traffic Volume Prediction

**Author:** Mahdi Farhani

**Email:** mm.farhani@gmail.com

This notebook explores the [Metro Interstate Traffic Volume](#) dataset. The aim is to perform exploratory data analysis (EDA) and establish a baseline regression model to predict the hourly traffic volume on Interstate 94, Minneapolis, using weather, holiday, and temporal features.

---

### 1.1 Learning Objectives

- Understand the structure and features of the traffic dataset.
- Perform exploratory data analysis and visualize key trends.
- Engineer relevant features for modeling.
- Build and evaluate a simple baseline regression model.

### 1.2 Dataset Introduction

The **Metro Interstate Traffic Volume** dataset provides hourly measurements of traffic volume along Interstate 94 in Minneapolis, Minnesota, USA, from 2012 to 2018. The data was collected by a traffic sensor located near a major interstate exit, and is closely linked to time, weather conditions, and public holidays.

This dataset is widely used for urban mobility research, traffic forecasting, and exploring the relationship between **weather phenomena** and **traffic patterns**. By analyzing the data, we can build predictive models to estimate road congestion, investigate trends over time (seasonal variations, rush hours, workdays vs holidays), and measure the impact of weather and calendar events on urban transportation.

- **Instances (Rows):** 48,204
- **Period:** 2012-2018 (hourly intervals)
- **Source:** [UCI Machine Learning Repository](#)

#### Goal:

Our goal is to predict the **traffic volume** (number of vehicles) passing an interstate segment per hour, using available features (weather, time, holiday status, etc.).

---

## 1.3 Column Descriptions

Column	Type	Description
holiday	Category	Name of the holiday, or “None” if not a public holiday
temp	Float	Hourly average temperature (Kelvin)
rain_1h	Float	Amount of rain in the preceding 1 hour (mm)
snow_1h	Float	Amount of snow in the preceding 1 hour (mm)
clouds_all	Integer	Percentage of sky covered by clouds
weather_main	Category	Main weather condition (e.g., Clear, Clouds, Rain, Snow, Mist)
weather_description	Category	More detailed weather condition (e.g., scattered clouds, light rain)
date_time	DateTime	Date and hour of the observation (local time)
traffic_volume	Integer	<b>Target Variable:</b> Number of vehicles (per hour)

## 1.4 1. Dataset Loading and Preview

Let’s load the Metro Interstate Traffic Volume dataset, specify column names, and display the first few rows to verify the data was imported correctly.

```
[769]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.preprocessing import StandardScaler
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LinearRegression
from sklearn.metrics import mean_absolute_error, mean_squared_error, r2_score

#plt setting
sns.set(style="whitegrid")
plt.rcParams["figure.figsize"] = (7, 4)

# Load the dataset
filename = '../data/Metro_Interstate_Traffic_Volume.csv'
dataframe= pd.read_csv(filename)
print (f"Dataset loaded successfully. {dataframe.shape[0]} rows and {dataframe.
↪shape[1]} columns.")
dataframe.head()
```

Dataset loaded successfully. 48204 rows and 9 columns.

```
[769]:  holiday    temp  rain_1h  snow_1h  clouds_all  weather_main  \
0      NaN    288.28      0.0      0.0          40         Clouds
1      NaN    289.36      0.0      0.0          75         Clouds
2      NaN    289.58      0.0      0.0          90         Clouds
```

3	NaN	290.13	0.0	0.0	90	Clouds
4	NaN	291.14	0.0	0.0	75	Clouds

	weather_description	date_time	traffic_volume
0	scattered clouds	2012-10-02 09:00:00	5545
1	broken clouds	2012-10-02 10:00:00	4516
2	overcast clouds	2012-10-02 11:00:00	4767
3	overcast clouds	2012-10-02 12:00:00	5026
4	broken clouds	2012-10-02 13:00:00	4918

### 1.4.1 Initial Data Exploration

In this section, we will: - Check for missing values in each column - Review the overall structure and data types - Display basic descriptive statistics for numerical columns

```
[770]: print("Missing values per column:")
print(dataframe.isnull().sum())
print("----")

print("General DataFrame Info:")
dataframe.info()
print("----")

print("Dataset Statistical Overview:")
print(dataframe.describe())
print("----")
```

```
Missing values per column:
holiday                48143
temp                   0
rain_1h                0
snow_1h                0
clouds_all             0
weather_main           0
weather_description    0
date_time              0
traffic_volume         0
dtype: int64
---
```

```
General DataFrame Info:
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 48204 entries, 0 to 48203
Data columns (total 9 columns):
#   Column                Non-Null Count  Dtype
---  -
0   holiday                61 non-null    object
1   temp                  48204 non-null float64
2   rain_1h               48204 non-null float64
```

```

3   snow_1h          48204 non-null float64
4   clouds_all       48204 non-null int64
5   weather_main     48204 non-null object
6   weather_description 48204 non-null object
7   date_time        48204 non-null object
8   traffic_volume   48204 non-null int64
dtypes: float64(3), int64(2), object(4)
memory usage: 3.3+ MB
---
```

Dataset Statistical Overview:

	temp	rain_1h	snow_1h	clouds_all	traffic_volume
count	48204.000000	48204.000000	48204.000000	48204.000000	48204.000000
mean	281.205870	0.334264	0.000222	49.362231	3259.818355
std	13.338232	44.789133	0.008168	39.015750	1986.860670
min	0.000000	0.000000	0.000000	0.000000	0.000000
25%	272.160000	0.000000	0.000000	1.000000	1193.000000
50%	282.450000	0.000000	0.000000	64.000000	3380.000000
75%	291.806000	0.000000	0.000000	90.000000	4933.000000
max	310.070000	9831.300000	0.510000	100.000000	7280.000000

---

**Feature Analysis** For each feature, we describe its type, value range, distribution, and possible impact on modeling. Use `describe()`, `unique()`, and `value_counts()` as needed. Pay attention to: - Outliers or odd values - Class imbalance for categorical features - Missing values - Data type conversion (integer/categorical)

```

[771]: print("Column types:\n")
        print(dataframe.dtypes)
        print("\nUnique columns:", dataframe.columns.tolist())

        print("Numerical Features Summary:")
        display(dataframe.describe().T)

        categorical_cols = ['holiday', 'weather_main', 'weather_description',
                             ↪ 'date_time']
        print("Categorical Features Distinct Values:")
        for col in categorical_cols:
            print(f"--- Column: {col}")
            print(dataframe[col].value_counts().head(10))
            print()

        numeric_cols = dataframe.select_dtypes(include=['float64', 'int64']).columns

        for col in numeric_cols:
```

```

print(f"---\nColumn: {col}")
print("Stats:")
print(dataframe[col].describe())
print("Nulls:", dataframe[col].isnull().sum())
print("Unique values:", dataframe[col].nunique())
print()

for col in categorical_cols:
    print(f"---\nColumn: {col}")
    print("Unique values/count:", dataframe[col].nunique())
    print("Sample values:", dataframe[col].unique()[:10])
    print("Top 10 most common:")
    print(dataframe[col].value_counts().head(10))
    print()

```

Column types:

```

holiday          object
temp             float64
rain_1h          float64
snow_1h          float64
clouds_all       int64
weather_main     object
weather_description object
date_time        object
traffic_volume   int64
dtype: object

```

Unique columns: ['holiday', 'temp', 'rain\_1h', 'snow\_1h', 'clouds\_all', 'weather\_main', 'weather\_description', 'date\_time', 'traffic\_volume']

Numerical Features Summary:

	count	mean	std	min	25%	50% \
temp	48204.0	281.205870	13.338232	0.0	272.16	282.45
rain_1h	48204.0	0.334264	44.789133	0.0	0.00	0.00
snow_1h	48204.0	0.000222	0.008168	0.0	0.00	0.00
clouds_all	48204.0	49.362231	39.015750	0.0	1.00	64.00
traffic_volume	48204.0	3259.818355	1986.860670	0.0	1193.00	3380.00

	75%	max
temp	291.806	310.07
rain_1h	0.000	9831.30
snow_1h	0.000	0.51
clouds_all	90.000	100.00
traffic_volume	4933.000	7280.00

Categorical Features Distinct Values:

```

--- Column: holiday
holiday

```

Labor Day	7
Thanksgiving Day	6
Christmas Day	6
New Years Day	6
Martin Luther King Jr Day	6
Columbus Day	5
Veterans Day	5
Washingtons Birthday	5
Memorial Day	5
Independence Day	5

Name: count, dtype: int64

--- Column: weather\_main

weather_main	
Clouds	15164
Clear	13391
Mist	5950
Rain	5672
Snow	2876
Drizzle	1821
Haze	1360
Thunderstorm	1034
Fog	912
Smoke	20

Name: count, dtype: int64

--- Column: weather\_description

weather_description	
sky is clear	11665
mist	5950
overcast clouds	5081
broken clouds	4666
scattered clouds	3461
light rain	3372
few clouds	1956
light snow	1946
Sky is Clear	1726
moderate rain	1664

Name: count, dtype: int64

--- Column: date\_time

date_time	
2013-05-19 10:00:00	6
2013-04-18 22:00:00	6
2012-12-16 21:00:00	5
2012-10-25 15:00:00	5
2013-12-16 10:00:00	5
2018-09-20 18:00:00	5

```
2016-11-18 14:00:00    5
2013-05-31 02:00:00    5
2012-11-11 04:00:00    5
2016-12-25 21:00:00    5
Name: count, dtype: int64
```

---

```
Column: temp
Stats:
count    48204.000000
mean      281.205870
std       13.338232
min        0.000000
25%       272.160000
50%       282.450000
75%       291.806000
max       310.070000
Name: temp, dtype: float64
Nulls: 0
Unique values: 5843
```

---

```
Column: rain_1h
Stats:
count    48204.000000
mean        0.334264
std       44.789133
min        0.000000
25%        0.000000
50%        0.000000
75%        0.000000
max      9831.300000
Name: rain_1h, dtype: float64
Nulls: 0
Unique values: 372
```

---

```
Column: snow_1h
Stats:
count    48204.000000
mean        0.000222
std        0.008168
min        0.000000
25%        0.000000
50%        0.000000
75%        0.000000
max        0.510000
Name: snow_1h, dtype: float64
```

Nulls: 0  
Unique values: 12

---

Column: clouds\_all

Stats:

count 48204.000000  
mean 49.362231  
std 39.015750  
min 0.000000  
25% 1.000000  
50% 64.000000  
75% 90.000000  
max 100.000000

Name: clouds\_all, dtype: float64

Nulls: 0

Unique values: 60

---

Column: traffic\_volume

Stats:

count 48204.000000  
mean 3259.818355  
std 1986.860670  
min 0.000000  
25% 1193.000000  
50% 3380.000000  
75% 4933.000000  
max 7280.000000

Name: traffic\_volume, dtype: float64

Nulls: 0

Unique values: 6704

---

Column: holiday

Unique values/count: 11

Sample values: [nan 'Columbus Day' 'Veterans Day' 'Thanksgiving Day' 'Christmas Day'

'New Years Day' 'Washingtons Birthday' 'Memorial Day' 'Independence Day'  
'State Fair']

Top 10 most common:

holiday	
Labor Day	7
Thanksgiving Day	6
Christmas Day	6
New Years Day	6
Martin Luther King Jr Day	6
Columbus Day	5



Veterans Day	5
Washingtons Birthday	5
Memorial Day	5
Independence Day	5

Name: count, dtype: int64

---

Column: weather\_main  
Unique values/count: 11  
Sample values: ['Clouds' 'Clear' 'Rain' 'Drizzle' 'Mist' 'Haze' 'Fog'  
'Thunderstorm'  
'Snow' 'Squall']  
Top 10 most common:  
weather\_main

Clouds	15164
Clear	13391
Mist	5950
Rain	5672
Snow	2876
Drizzle	1821
Haze	1360
Thunderstorm	1034
Fog	912
Smoke	20

Name: count, dtype: int64

---

Column: weather\_description  
Unique values/count: 38  
Sample values: ['scattered clouds' 'broken clouds' 'overcast clouds' 'sky is  
clear'  
'few clouds' 'light rain' 'light intensity drizzle' 'mist' 'haze' 'fog']  
Top 10 most common:  
weather\_description

sky is clear	11665
mist	5950
overcast clouds	5081
broken clouds	4666
scattered clouds	3461
light rain	3372
few clouds	1956
light snow	1946
Sky is Clear	1726
moderate rain	1664

Name: count, dtype: int64

---

Column: date\_time

```

Unique values/count: 40575
Sample values: ['2012-10-02 09:00:00' '2012-10-02 10:00:00' '2012-10-02
11:00:00'
'2012-10-02 12:00:00' '2012-10-02 13:00:00' '2012-10-02 14:00:00'
'2012-10-02 15:00:00' '2012-10-02 16:00:00' '2012-10-02 17:00:00'
'2012-10-02 18:00:00']
Top 10 most common:
date_time
2013-05-19 10:00:00    6
2013-04-18 22:00:00    6
2012-12-16 21:00:00    5
2012-10-25 15:00:00    5
2013-12-16 10:00:00    5
2018-09-20 18:00:00    5
2016-11-18 14:00:00    5
2013-05-31 02:00:00    5
2012-11-11 04:00:00    5
2016-12-25 21:00:00    5
Name: count, dtype: int64

```

#### 1.4.2 Notes:

- All feature statistics, unique values, and data types were explored.
- Missing values—and if they exist—should be imputed or dropped before modeling.
- Extracting time-based features from `date_time` helps reveal daily/weekly patterns and will be used in feature engineering.
- Always visualize your target and key features to catch outliers and distribution properties.

```

[772]: dataframe['date_time'] = pd.to_datetime(dataframe['date_time'])

dataframe['hour'] = dataframe['date_time'].dt.hour
dataframe['weekday'] = dataframe['date_time'].dt.weekday
dataframe['month'] = dataframe['date_time'].dt.month
dataframe['year'] = dataframe['date_time'].dt.year

print("Sample extracted time columns:")
display(dataframe[['date_time', 'hour', 'weekday', 'month', 'year']].head())

dataframe['temp'] = dataframe['temp'] - 273.15
display(dataframe[['temp']].head())

dataframe['temp'].hist(bins=30)
plt.xlabel('Temperature (°C)')
plt.ylabel('Frequency')
plt.title('Distribution of Temperature (Celsius)')
plt.show()

```

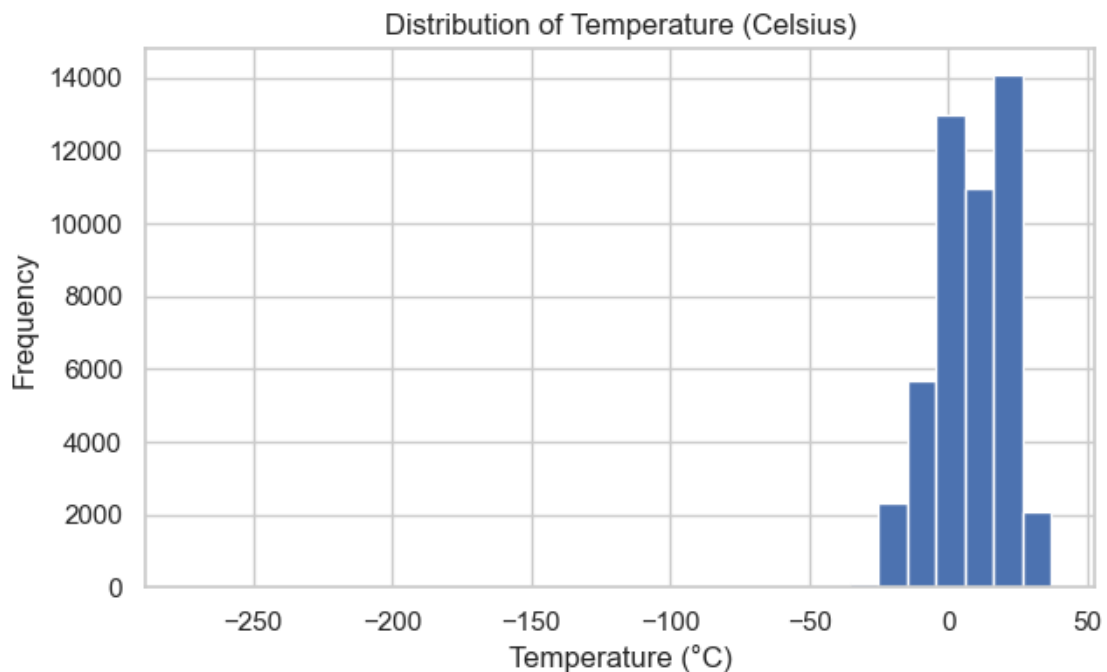
```
dataframe[dataframe['temp'] < -25].shape

dataframe = dataframe[dataframe['temp'] >= -25]
```

Sample extracted time columns:

	date_time	hour	weekday	month	year
0	2012-10-02 09:00:00	9	1	10	2012
1	2012-10-02 10:00:00	10	1	10	2012
2	2012-10-02 11:00:00	11	1	10	2012
3	2012-10-02 12:00:00	12	1	10	2012
4	2012-10-02 13:00:00	13	1	10	2012

	temp
0	15.13
1	16.21
2	16.43
3	16.98
4	17.99



**Visualizing the Distribution of Traffic Volume** The following code creates a histogram of the target variable `traffic_volume` using Seaborn. The `bins=40` parameter ensures that the range of traffic volumes is divided into 40 intervals, and `kde=True` overlays a Kernel Density Estimate curve to show the probability distribution.


- **Purpose:**

To understand the range, central tendency, spread, skewness, and possible outliers in the hourly traffic counts.

- **Interpretation:**

If the histogram is right-skewed (a long tail on the right), it means most hours have lower to moderate traffic, but there are some hours with exceptionally high traffic.

```
[773]: plt.figure()
sns.histplot(dataframe['traffic_volume'], bins=40, kde=True)
plt.title("Distribution of Traffic Volume")
plt.xlabel("Traffic Volume")
plt.ylabel("Frequency")
plt.show()
```




### 1.4.3 Distribution Analysis of Traffic Volume

- The histogram and KDE curve reveal a **multi-modal distribution** of hourly traffic volume.
- **Multiple peaks** suggest distinct traffic patterns corresponding to different times of the day (e.g., low traffic at night, peaks during rush hours).
- The **first mode** (0–1000) likely represents nighttime or early hours; the **second and third modes** (around 3000 and 5000) are probably linked to midday and evening peaks.
- The long right-tail and rare high values (>7000) may indicate outliers.
- **Conclusion:**

This distribution suggests the importance of time-based feature engineering. Further, separating analysis by hour-of-day may reveal even clearer traffic patterns.

```
[774]: hourly_traffic = dataframe.groupby('hour')['traffic_volume'].mean()

plt.figure()
sns.barplot(x=hourly_traffic.index, y=hourly_traffic.values,
            palette="crest", legend=False, hue=0.1)
plt.title("Average Traffic Volume by Hour of Day")
plt.xlabel("Hour of Day (0-23)")
plt.ylabel("Average Traffic Volume")
plt.show()
```



#### 1.4.4 Traffic Volume Patterns: Overall vs. Hourly

The overall distribution of `traffic_volume` is multi-modal, indicating the presence of fundamentally different traffic behaviors throughout a 24-hour cycle— from quiet nighttime hours to intense rush hour peaks.

By examining the **average traffic volume for each hour**, the sources of the main peaks and valleys in the overall histogram become clear. There is a sharp increase in traffic starting around 6:00, with sustained high levels during business hours and particularly pronounced peaks in the late afternoon. The overnight hours (0–5 AM) consistently show low traffic volumes.

##### Key modeling insight:

Rather than treating traffic volume as a single homogeneous variable, it is essential to incorporate

time-based features when building predictive models. This temporal breakdown explains the multi-modal distribution and should inform future feature engineering and data preprocessing steps.

### Additional EDA Recommendations

- **Visualize outliers and extreme values** to prevent their undue influence on downstream modeling.
- **Investigate temporal patterns** not only by hour but also by day of week and holidays, as these may reveal further structure.
- **Carefully check for missing or anomalous data** to ensure data quality.
- **Explore relationships** between traffic volume and other explanatory features (e.g., weather, temperature) using correlation plots and visual explorations.
- **Document any data biases or coverage gaps** to contextualize your findings and model limitations.

#### 1.4.5 Outlier Detection in Traffic Volume Data

We identified outliers in the target variable `traffic_volume` using boxplots and the IQR statistical method. These extreme values can be caused by unusual events such as road closures, accidents, or data recording errors. For further modeling, we review and decide whether to exclude, cap, or keep these records based on their context and distribution.

**Key points:** - Outliers were detected both in very low and extremely high traffic volumes. - We also screened for impossible values (negative or zero traffic volumes). - Visualization by hour highlighted that certain outliers happen typically during off-peak hours.

**Boxplot Analysis of Traffic Volume** The boxplot of `traffic_volume` shows the distribution and presence of outliers in the dataset. Key observations:

- **Interquartile Range (IQR):** Most data points are centered between approximately 1,200 and 4,800 vehicles per hour (the edges of the box).
- **Median:** The median (central black line) is close to the center of the IQR, indicating a relatively symmetric core distribution for the main traffic flow.
- **Whiskers:** The whiskers extend to nearly 0 and just above 7,000. These set the boundaries for typical traffic volumes, beyond which any data point is considered an outlier by the classic boxplot definition.
- **Outliers:** Only a handful of data points are considered outliers (shown as dots beyond the whiskers). The majority of `traffic_volume` values are inside the normal range.
- **No significant skewness in box:** The box itself is not extremely unbalanced to either side, supporting our previous observation that the main distribution is not highly skewed but may be multi-modal due to rush hours.

*Practical note:*

A small number of records with extremely low (close to zero) or extremely high (above 7,000) traffic volumes exist, warranting additional review for possible data quality issues or special events (e.g., maintenance, road closure).

**Hourly Boxplot Analysis of Traffic Volume** The hourly boxplots of `traffic_volume` reveal critical patterns about data variability, outliers, and commuter behavior throughout the day:

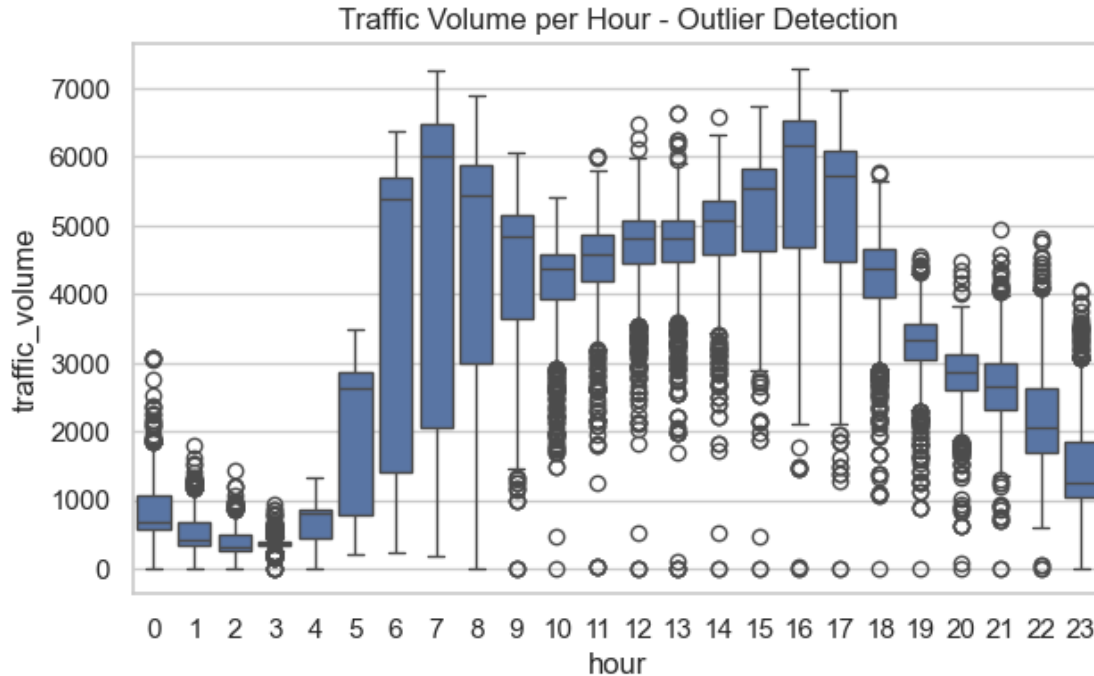
- **Night Hours (0-5):**
  - Traffic volumes are low with little variation.
  - Outliers (black circles) are much more common, many of which are near zero—possibly due to near-empty roads or special conditions (road closure, sensor anomalies).
- **Morning Peak (6-9):**
  - Sharp increase in both the median and spread of traffic volumes, corresponding to the morning commute.
  - The variability within these hours is high, and a wider range of outliers can be observed.
- **Midday (10-15):**
  - Medians are lower and variability decreases compared to rush hours.
  - Still, outliers exist at both the low and high traffic ranges, reflecting variable midday traffic patterns.
- **Evening Peak (16-18):**
  - Another distinct increase in the median and spread is seen.
  - The range is broad, with significant outliers mainly on the lower side (possibly sudden drops in traffic due to incidents).
- **Late Evening (19-23):**
  - Median volumes drop and spread narrows.
  - Outliers increase again, especially at low traffic volumes.

**Observations:** - Traffic volume outliers are not distributed equally across hours—they are much more likely at night and during less busy midday hours. - The multi-modal, hour-dependent pattern shows why a single overall `traffic_volume` boxplot lacks nuance. - Data with near-zero volumes at unusual hours may warrant closer inspection for error or special events.

*Practical tip:*

When cleaning or modeling, consider outlier behavior on an hourly basis, not just across the full dataset. Removing outliers only globally could lead to discarding valuable nighttime or weekend observations.

```
[775]: plt.figure(figsize=(7,4))
sns.boxplot(x=dataframe['hour'], y=dataframe['traffic_volume'])
plt.title('Traffic Volume per Hour - Outlier Detection')
plt.show()
```



**Interpretation of the Hourly Traffic Volume Boxplot** This hour-by-hour boxplot offers deep insight into the temporal dynamics and variability of traffic volume:

- **1. Daily Traffic Pattern:**
  - During late night and early morning hours (0-5), traffic is consistently low, reflecting minimal nocturnal driving.
  - Starting from 6 AM, a sharp rise is observed, peaking between 7 and 9 AM, corresponding to morning rush hour.
  - After 9 AM, traffic volume decreases and remains relatively steady throughout midday.
  - A second clear peak appears around 4-6 PM (16-18), representing evening rush hour and commuters returning home.
  - After 7 PM, volumes gradually decline again towards nighttime levels.
- **2. Outlier Distribution:**
  - Outliers are heavily concentrated during nighttime and early morning hours, often at or near zero—possibly indicating empty roads, maintenance events, or sensor issues.
  - Both morning and evening peak hours show a broad interquartile range and a significant number of outliers, reflecting variable, sometimes extreme congestion or unusual traffic drops (accidents, sudden events).
  - Midday and night hours feature fewer outliers, mostly at the lower end, possibly due to sporadic dips in traffic.
- **3. Spread and Variability:**
  - The width of the box (interquartile range) is greatest during rush hours, indicating higher variability in commuter behavior.
  - In quiet nighttime hours, the box is narrow, signaling highly consistent (low) traffic volumes.



- **4. Importance of Temporal Analysis:**

- Traffic behavior and outlier occurrence are highly time-dependent. Overall (global) analyses may miss crucial details present only at specific times.
- Thus, data cleaning, outlier handling, or even model training should incorporate hourly segmentation to preserve important patterns and avoid loss of valuable edge-case data.

**In summary:**

This visualization confirms that only time-aware analysis can accurately capture the real-world behavior of traffic volume. Any data-cleansing or modeling effort should respect the strong temporal regularities and outlier characteristics revealed by the hourly plots.

In addition to hourly patterns, we investigated how traffic volume varies by day of week and holidays.


- **Weekdays vs. Weekends:** Average and distribution of traffic volume are distinctly different; weekends typically show lower peaks and more uniform volumes throughout the day.
- **Day of Week:** Each day exhibits a unique profile, with Mondays and Fridays sometimes showing higher variability.
- **Holidays:** Traffic volumes during public holidays are usually much lower (unless coinciding with special events). The distribution is noticeably distinct from regular weekdays.
- **Combined Effects:** A heatmap of average traffic by hour and day of week reveals complex, structured patterns—e.g., rush hours are much less pronounced on weekends and holidays.

**Practical insight:**

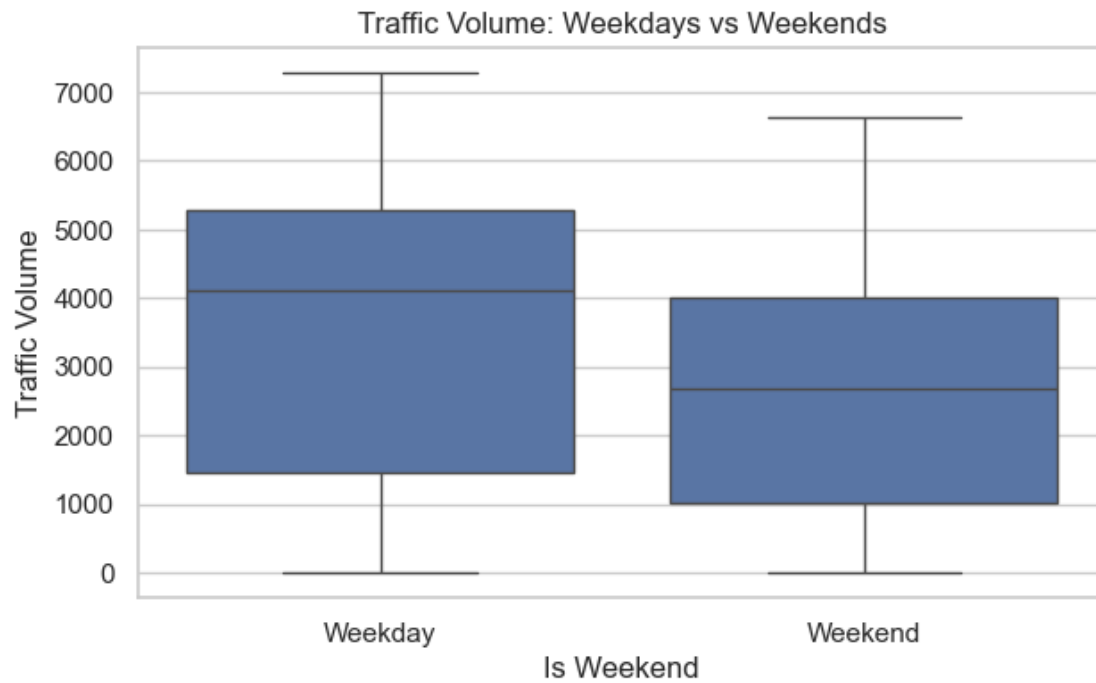
Time-based features, especially day-of-week and holiday flags, can substantially improve the predictive power of traffic forecasting models.

```
[776]: dataframe['is_weekend'] = dataframe['weekday'] >= 5

plt.figure(figsize=(7,4))
sns.barplot(x='weekday', y='traffic_volume', data=dataframe, errorbar=None)
plt.xlabel('Day of Week (0=Mon)')
plt.ylabel('Average Traffic Volume')
plt.title('Average Traffic Volume by Day of Week')
plt.show()
```

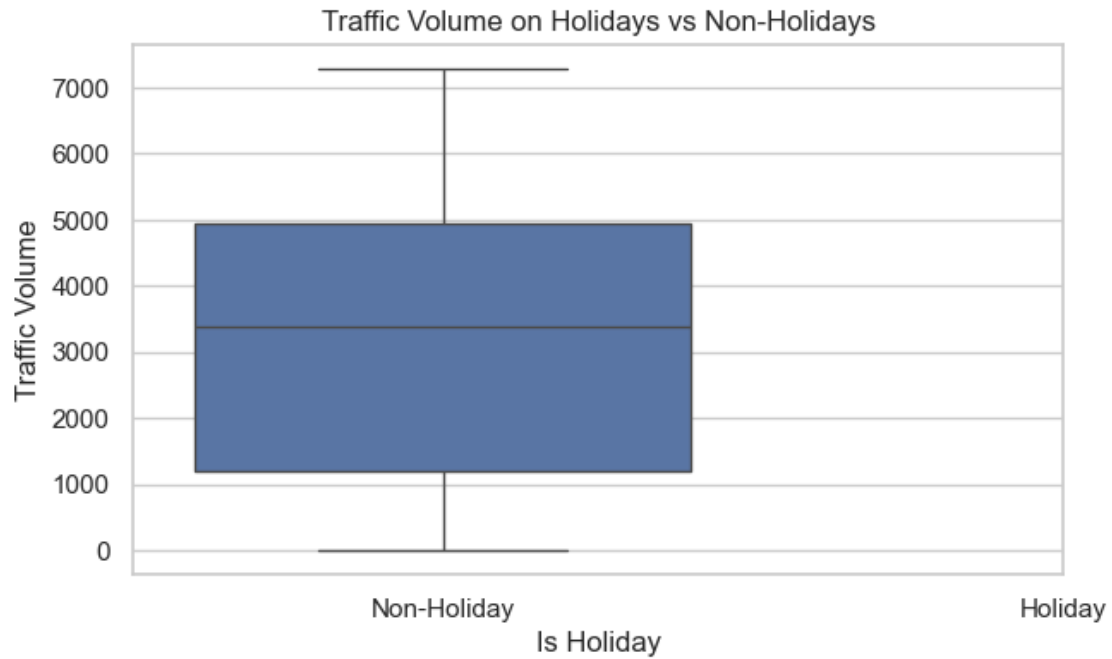


```
[777]: plt.figure(figsize=(7,4))
sns.boxplot(x='is_weekend', y='traffic_volume', data=dataframe)
plt.xlabel('Is Weekend')
plt.ylabel('Traffic Volume')
plt.title('Traffic Volume: Weekdays vs Weekends')
plt.xticks([0, 1], ['Weekday', 'Weekend'])
plt.show()
```




```
[778]: dataframe['is_holiday'] = dataframe['holiday'] != 'None'

plt.figure(figsize=(7,4))
sns.boxplot(x='is_holiday', y='traffic_volume', data=dataframe)
plt.xlabel('Is Holiday')
plt.ylabel('Traffic Volume')
plt.title('Traffic Volume on Holidays vs Non-Holidays')
plt.xticks([0, 1], ['Non-Holiday', 'Holiday'])
plt.show()
```



```
[779]: pivot = dataframe.pivot_table(values='traffic_volume',
                                     index='hour', columns='weekday', aggfunc='mean')

plt.figure(figsize=(12,6))
sns.heatmap(pivot, annot=True, fmt='.0f', cmap='YlGnBu')
plt.xlabel('Day of Week (0=Mon)')
plt.ylabel('Hour of Day')
plt.title('Mean Traffic Volume by Hour and Day of Week')
plt.show()
```



## 1.5 Temporal Traffic Pattern Analysis

### 1.5.1 1. Average Traffic Volume by Day of Week

- **Observation:** Traffic peaks from Monday to Friday (especially midweek), then drops sharply on weekends (Saturday and Sunday).
- **Insight:** A strong weekly pattern is evident; weekdays show higher, more sustained traffic, while weekends are noticeably lighter and smoother.

### 1.5.2 2. Traffic Volume: Weekdays vs. Weekends

- **Observation:** Weekdays have a higher average and wider spread of traffic volume (boxplot), indicating strong commuter influence.
- **Insight:** Weekend traffic is both lower and less variable, pointing to a fundamentally different travel behavior, likely less commuting and more recreational trips.

### 1.5.3 3. Traffic Volume on Holidays vs. Non-Holidays

- **Observation:** On public holidays, both average and spread of traffic decrease significantly.
- **Insight:** Holidays suppress urban traffic demand, suggesting many people avoid work and routine travel.

#### 1.5.4 4. Mean Traffic Volume by Hour and Day of Week (Heatmap)

- **Observation:** Pronounced morning and evening rush hours on weekdays (7–9 AM, 4–6 PM) are greatest, while on weekends and holidays these peaks flatten or nearly disappear.
- **Insight:** Both the hour of day, day of week, and holiday status are crucial for modeling real-world traffic dynamics.

---

#### Conclusion:

Traffic is not only a function of the clock, but also depends deeply on the day and holiday effects. Including these temporal features can dramatically improve the predictive performance of machine learning models.


---

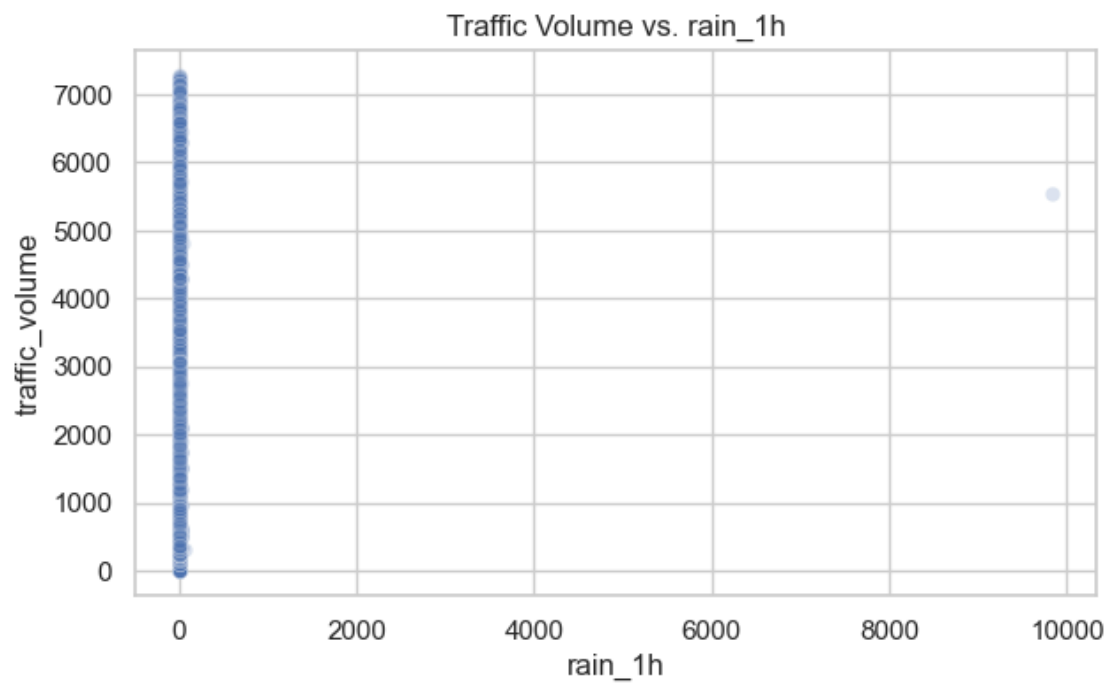
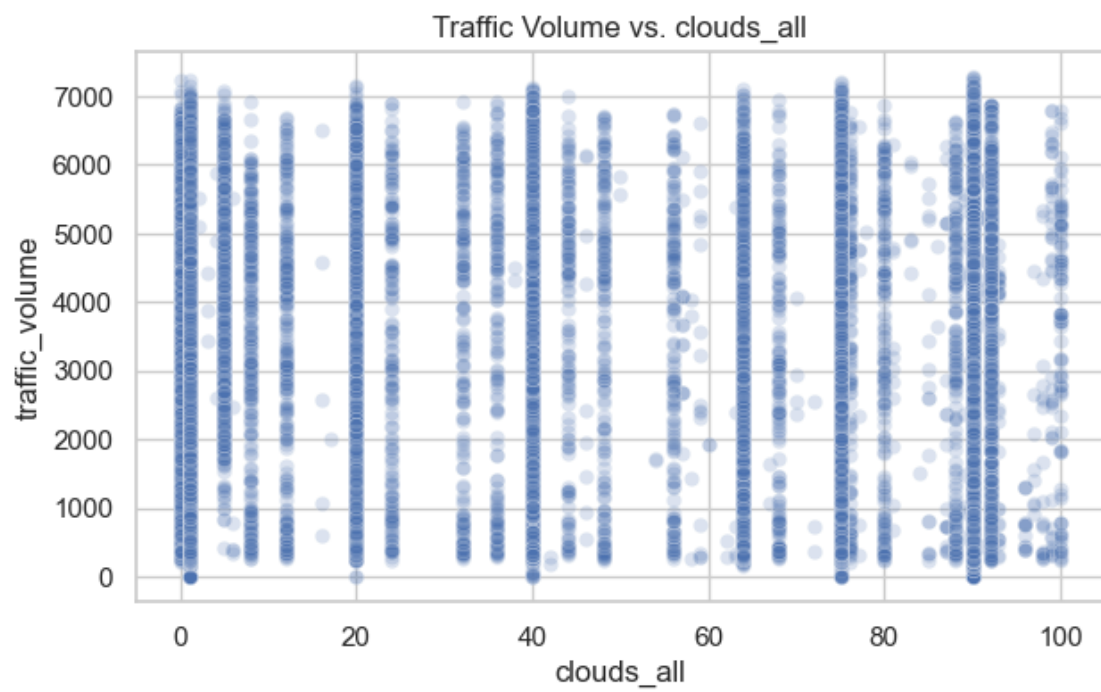
### 1.6 Exploring Relationships: Traffic Volume & Explanatory Features

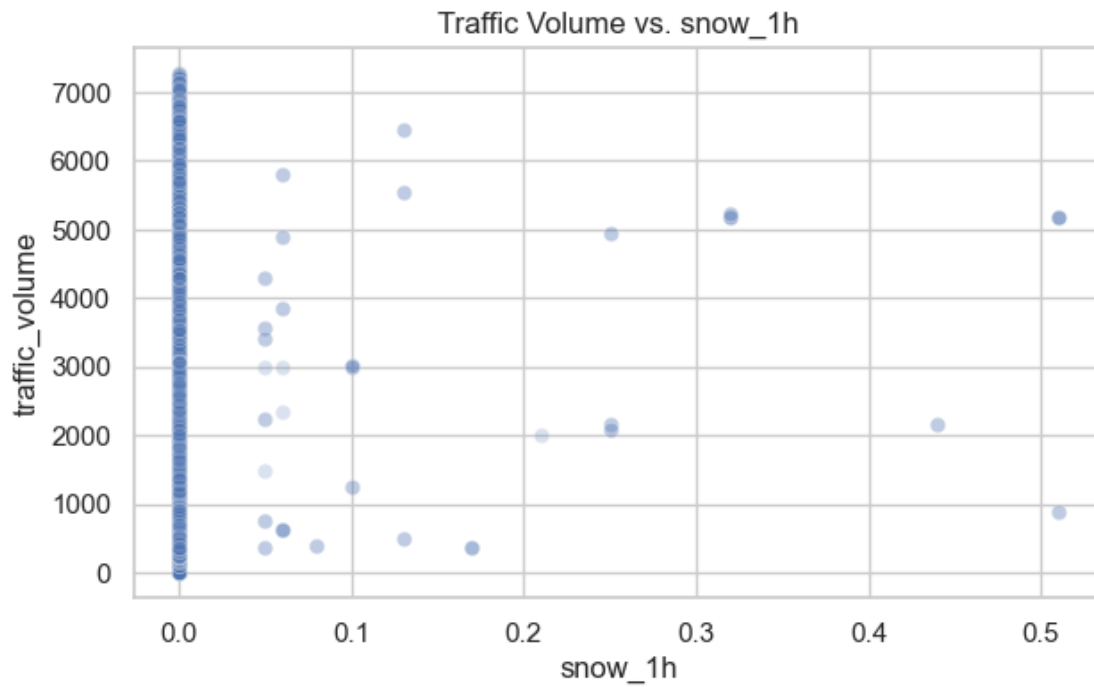
#### 1.6.1 1. Correlation Analysis

Use correlation coefficients to quantify the linear relationship between `traffic_volume` and numeric features (e.g., `temp`, `rain_1h`, `clouds_all`):

```
[780]: features = ['temp', 'clouds_all', 'rain_1h', 'snow_1h']
for col in features:
    plt.figure(figsize=(7,4))
    sns.scatterplot(x=dataframe[col], y=dataframe['traffic_volume'], alpha=0.2)
    plt.title(f'Traffic Volume vs. {col}')
    plt.show()
```








```
[781]: plt.figure(figsize=(7,4))
sns.boxplot(x='weather_main', y='traffic_volume', data=dataframe)
plt.xticks(rotation=45)
plt.title('Traffic Volume by Weather Condition')
plt.show()
```





## 1.6.2 Data Cleaning and Feature Engineering

**1 Outlier Removal in Rain Data** To eliminate unrealistic outlier values, all records where `rain_1h` is greater than or equal to 500 (most likely data entry errors) were removed. #####

**2. Binary Features for Rain and Snow** Two new binary features were engineered to indicate the presence of rain or snow for each record. #####

**3. Weather Condition Text Standardization** String values in the `weather_main` column were standardized by stripping extra spaces and converting all entries to lowercase. #####

**4. Grouping Rare Weather Categories** Weather conditions that appear less than 100 times in the dataset were grouped into a single “Other” category as part of cleaning and simplifying the feature.

```
[782]: dataframe = dataframe[dataframe['rain_1h'] < 500]

dataframe['is_rain'] = (dataframe['rain_1h'] > 0).astype(int)
dataframe['is_snow'] = (dataframe['snow_1h'] > 0).astype(int)

dataframe['weather_main'] = dataframe['weather_main'].str.strip().str.lower()
rare_weather = dataframe['weather_main'].
    ↪value_counts()[dataframe['weather_main'].value_counts() < 100].index
dataframe['weather_main_cleaned'] = dataframe['weather_main'].
    ↪replace(rare_weather, 'other')
```

### 1.6.3 Handling Categorical Features: One-Hot Encoding for weather\_main

To prepare the `weather_main` categorical feature for machine learning models, we applied **one-hot encoding**. This process converts each unique value in the `weather_main` column (except the first, due to `drop_first=True`) into a separate binary column, indicating the presence (1) or absence (0) of that weather type for each observation. This step is critical to ensure that machine learning algorithms properly handle categorical variables without inferring any ordinal relationship between them.

Note: After this step, the original `weather_main` column is replaced by several new columns (e.g., `weather_main_clouds`, `weather_main_rain`, etc.), each representing one of the weather conditions.

```
[783]: dataframe = pd.get_dummies(dataframe, columns=['weather_main_cleaned'])
dataframe.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Index: 48075 entries, 0 to 48203
Data columns (total 27 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   holiday                               61 non-null     object
1   temp                                  48075 non-null  float64
2   rain_1h                               48075 non-null  float64
3   snow_1h                               48075 non-null  float64
4   clouds_all                            48075 non-null  int64
5   weather_main                          48075 non-null  object
6   weather_description                   48075 non-null  object
7   date_time                             48075 non-null  datetime64[ns]
8   traffic_volume                        48075 non-null  int64
9   hour                                  48075 non-null  int32
10  weekday                               48075 non-null  int32
11  month                                 48075 non-null  int32
12  year                                  48075 non-null  int32
13  is_weekend                           48075 non-null  bool
14  is_holiday                           48075 non-null  bool
15  is_rain                               48075 non-null  int32
16  is_snow                               48075 non-null  int32
17  weather_main_cleaned_clear            48075 non-null  bool
18  weather_main_cleaned_clouds           48075 non-null  bool
19  weather_main_cleaned_drizzle          48075 non-null  bool
20  weather_main_cleaned_fog              48075 non-null  bool
21  weather_main_cleaned_haze             48075 non-null  bool
22  weather_main_cleaned_mist             48075 non-null  bool
23  weather_main_cleaned_other            48075 non-null  bool
24  weather_main_cleaned_rain             48075 non-null  bool
25  weather_main_cleaned_snow             48075 non-null  bool
26  weather_main_cleaned_thunderstorm     48075 non-null  bool
dtypes: bool(12), datetime64[ns](1), float64(3), int32(6), int64(2), object(3)
memory usage: 5.3+ MB
```

### 1.6.4 Dropping Unnecessary Columns Before Train/Test Split

To prevent information leakage and ensure that only relevant features are used for model training and evaluation, it is essential to drop unnecessary columns before splitting the data into training and test sets. These may include original categorical columns that have been replaced by one-hot encoding, columns not related to modeling (such as IDs or timestamps), or any features not intended for modeling.

For example, after one-hot encoding `weather_main`, the original column can be safely dropped (if not already by `get_dummies`), and similarly, irrelevant columns like `date_time` or identifiers can be excluded:

```
[784]: selected_columns = [ 'is_rain', 'is_snow',
                           'hour', 'weekday',
                           'is_weekend', 'is_holiday',
                           'weather_main_cleaned_other', 'weather_main_cleaned_clear', 'weather_main_cleaned_clouds',
                           'weather_main_cleaned_drizzle',
                           'weather_main_cleaned_fog', 'weather_main_cleaned_haze',
                           'weather_main_cleaned_mist',
                           'weather_main_cleaned_rain', 'weather_main_cleaned_snow',
                           'weather_main_cleaned_thunderstorm', 'traffic_volume']

new_dataframe=dataframe[selected_columns].copy()
new_dataframe.info()
```

```
<class 'pandas.core.frame.DataFrame'>
```

```
Index: 48075 entries, 0 to 48203
```

```
Data columns (total 17 columns):
```

#	Column	Non-Null Count	Dtype
0	is_rain	48075 non-null	int32
1	is_snow	48075 non-null	int32
2	hour	48075 non-null	int32
3	weekday	48075 non-null	int32
4	is_weekend	48075 non-null	bool
5	is_holiday	48075 non-null	bool
6	weather_main_cleaned_other	48075 non-null	bool
7	weather_main_cleaned_clear	48075 non-null	bool
8	weather_main_cleaned_clouds	48075 non-null	bool
9	weather_main_cleaned_drizzle	48075 non-null	bool
10	weather_main_cleaned_fog	48075 non-null	bool
11	weather_main_cleaned_haze	48075 non-null	bool
12	weather_main_cleaned_mist	48075 non-null	bool
13	weather_main_cleaned_rain	48075 non-null	bool
14	weather_main_cleaned_snow	48075 non-null	bool
15	weather_main_cleaned_thunderstorm	48075 non-null	bool
16	traffic_volume	48075 non-null	int64

```
dtypes: bool(12), int32(4), int64(1)
```

memory usage: 2.0 MB

### 1.6.5 Train/Test Split

To evaluate the generalization performance of our model, we split the dataset into training and testing sets.

Here, 80% of the data is used for training, and 20% is held out for testing.

```
[785]: X = new_dataframe.drop('traffic_volume', axis=1)
y = new_dataframe['traffic_volume']
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2,
↳ random_state=42)
```

### 1.6.6 Feature Scaling with StandardScaler

Many machine learning algorithms perform better when numerical features are standardized.

Here, we use `StandardScaler` to scale the features so that each has a mean of 0 and a standard deviation of 1.

```
[786]: scaler = StandardScaler()
X_train_scaled = scaler.fit_transform(X_train)
X_test_scaled = scaler.transform(X_test)
```

### 1.6.7 Baseline Modeling: Linear Regression

For our initial modeling step, we use a simple **Linear Regression** as a baseline.

This allows us to set a reference point for future, more advanced models.

**Steps:** 1. Instantiate a Linear Regression model. 2. Fit the model on the scaled training data. 3. Predict the target for the test set. 4. Calculate and report evaluation metrics: MAE, RMSE, and  $R^2$ .

```
[787]: baseline_model = LinearRegression()
baseline_model.fit(X_train_scaled, y_train)
y_pred = baseline_model.predict(X_test_scaled)
mae = mean_absolute_error(y_test, y_pred)
rmse = np.sqrt(mean_squared_error(y_test, y_pred))
r2 = r2_score(y_test, y_pred)

print(f"MAE: {mae:.2f}")
print(f"RMSE: {rmse:.2f}")
print(f"R²: {r2:.3f}")
```

MAE: 1582.63

RMSE: 1797.86

$R^2$ : 0.167

```
[788]: residuals = y_test - y_pred

plt.figure(figsize=(16, 4))
```

```


# 1. Actual vs Predicted
plt.subplot(1, 3, 1)
sns.scatterplot(x=y_test, y=y_pred, alpha=0.6)
plt.xlabel('Actual Traffic Volume')
plt.ylabel('Predicted Traffic Volume')
plt.title('Actual vs Predicted')
plt.plot([y_test.min(), y_test.max()], [y_test.min(), y_test.max()], '--r')
plt.grid(True)

# 2. Residual Plot
plt.subplot(1, 3, 2)
sns.scatterplot(x=y_pred, y=residuals, alpha=0.6)
plt.axhline(0, color='red', ls='--')
plt.xlabel('Predicted Traffic Volume')
plt.ylabel('Residuals')
plt.title('Residuals vs Predicted')
plt.grid(True)

# 3. Histogram of Residuals
plt.subplot(1, 3, 3)
sns.histplot(residuals, bins=30, kde=True, color='orange')
plt.xlabel('Residuals')
plt.title('Distribution of Residuals')
plt.grid(True)

plt.tight_layout()
plt.show()

```



## 1.7 Feature Engineering: Cyclical Encoding of Hour and Weekday

### 1.7.1 Overview

Many time-based features such as **hour** of the day and **weekday** are **cyclical** in nature:

- After hour 23 comes hour 0 (midnight resets the cycle)
- After Sunday (weekday=6) comes Monday (weekday=0)

Directly using integer values for these features in a linear model assumes a straight-line relationship,

which creates artificial discontinuities (e.g., difference between hour 23 and 0 is interpreted as “23” instead of a small circular gap).

**To overcome this** and preserve the cyclical relationships

```
[789]: new_dataframe['hour_sin'] = np.sin(2 * np.pi * new_dataframe['hour']/24)
new_dataframe['hour_cos'] = np.cos(2 * np.pi * new_dataframe['hour']/24)
new_dataframe['weekday_sin'] = np.sin(2 * np.pi * new_dataframe['weekday']/7)
new_dataframe['weekday_cos'] = np.cos(2 * np.pi * new_dataframe['weekday']/7)
new_dataframe.drop(['weekday', 'hour'], axis=1, inplace=True)

new_dataframe.info()

X = new_dataframe.drop('traffic_volume', axis=1)
y = new_dataframe['traffic_volume']
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2,
    random_state=42)

baseline_model = LinearRegression()
baseline_model.fit(X_train_scaled, y_train)
y_pred = baseline_model.predict(X_test_scaled)
mae = mean_absolute_error(y_test, y_pred)
rmse = np.sqrt(mean_squared_error(y_test, y_pred))
r2 = r2_score(y_test, y_pred)

print(f"MAE: {mae:.2f}")
print(f"RMSE: {rmse:.2f}")
print(f"R²: {r2:.3f}")

residuals = y_test - y_pred

plt.figure(figsize=(16, 4))

# 1. Actual vs Predicted
plt.subplot(1, 3, 1)
sns.scatterplot(x=y_test, y=y_pred, alpha=0.6)
plt.xlabel('Actual Traffic Volume')
plt.ylabel('Predicted Traffic Volume')
plt.title('Actual vs Predicted')
plt.plot([y_test.min(), y_test.max()], [y_test.min(), y_test.max()], '--r')
plt.grid(True)

# 2. Residual Plot
plt.subplot(1, 3, 2)
sns.scatterplot(x=y_pred, y=residuals, alpha=0.6)
plt.axhline(0, color='red', ls='--')
plt.xlabel('Predicted Traffic Volume')
plt.ylabel('Residuals')
```

```
plt.title('Residuals vs Predicted')
plt.grid(True)

# 3. Histogram of Residuals
plt.subplot(1, 3, 3)
sns.histplot(residuals, bins=30, kde=True, color='orange')
plt.xlabel('Residuals')
plt.title('Distribution of Residuals')
plt.grid(True)

plt.tight_layout()
plt.show()
```

```
<class 'pandas.core.frame.DataFrame'>
```

```
Index: 48075 entries, 0 to 48203
```

```
Data columns (total 19 columns):
```

#	Column	Non-Null Count	Dtype
0	is_rain	48075 non-null	int32
1	is_snow	48075 non-null	int32
2	is_weekend	48075 non-null	bool
3	is_holiday	48075 non-null	bool
4	weather_main_cleaned_other	48075 non-null	bool
5	weather_main_cleaned_clear	48075 non-null	bool
6	weather_main_cleaned_clouds	48075 non-null	bool
7	weather_main_cleaned_drizzle	48075 non-null	bool
8	weather_main_cleaned_fog	48075 non-null	bool
9	weather_main_cleaned_haze	48075 non-null	bool
10	weather_main_cleaned_mist	48075 non-null	bool
11	weather_main_cleaned_rain	48075 non-null	bool
12	weather_main_cleaned_snow	48075 non-null	bool
13	weather_main_cleaned_thunderstorm	48075 non-null	bool
14	traffic_volume	48075 non-null	int64
15	hour_sin	48075 non-null	float64
16	hour_cos	48075 non-null	float64
17	weekday_sin	48075 non-null	float64
18	weekday_cos	48075 non-null	float64


```
dtypes: bool(12), float64(4), int32(2), int64(1)
```

```
memory usage: 3.1 MB
```

```
MAE: 1582.63
```

```
RMSE: 1797.86
```

```
R2: 0.167
```



## Effect of Sine/Cosine Encoding for Cyclical Time Features

- **Observation:** Adding cyclical features (hour/weekday with sine/cosine encoding) did not meaningfully improve linear model fit.
- **Interpretation:** While sin/cos encoding allows the model to “see” the periodicity of time, a linear regression still cannot represent sharp peaks, plateaus, or step changes in traffic that arise in urban data (i.e., transitions at rush hour, weekends, holidays).
- **Residual Patterns:** The clear bifurcation and strong non-random structure in the residuals demonstrate the presence of distinct sub-populations (e.g., rush hour vs. off-peak), which linear models cannot capture without explicit interaction features or nonlinearity.
- **Takeaway:** Sine/cosine encoding is necessary for cyclical variables, but **not sufficient** for representing truly complex, non-linear relationships—especially when the true relationship with the target is not a simple sine wave!
- **Recommendation:** For further improvement, use tree-based or ensemble models (RandomForest, XGBoost) or explicitly add interaction features.

### 1.7.2 Observations from Linear Regression

Despite careful feature engineering—including cyclical encoding of **hour** and **weekday** using sin/cos—the linear regression model still struggles with the complexity of urban traffic volume prediction:

- There is significant **underfitting**, as demonstrated by:
  - Actual vs. Predicted plots with points clustering off the diagonal line or within tight bands.
  - Residual analysis showing structured errors (branches or multi-modal distributions).
- The error distribution (residuals) is often multi-peaked or wide, indicating the presence of **hidden sub-patterns** or seasonal effects that linear models cannot represent.
- Even proper encoding of cyclical features is **necessary but not sufficient**—relationships between variables and the target (traffic volume) are highly non-linear and involve complex interactions.

### 1.7.3 Why Move to Non-linear Models?

Tree-based ensemble models such as **RandomForest** can:

- Capture complex, non-linear interactions between features.
- Automatically handle feature splits, interactions, and variable importance.



- Produce significantly better performance and residual behavior in real-world, multi-modal data scenarios (such as traffic flow).

---

#### 1.7.4 Next Step

In the following step, we will train and evaluate a `RandomForestRegressor` on the same set of features. We expect the model to better capture the non-linear and multi-patterned nature of traffic volume, resulting in improved prediction metrics and more random (less structured) residuals.

---

```
[790]: from sklearn.ensemble import RandomForestRegressor

rf = RandomForestRegressor(n_estimators=100, random_state=0, n_jobs=-1)
rf.fit(X_train, y_train)

y_pred_rf = rf.predict(X_test)

mae_rf = mean_absolute_error(y_test, y_pred_rf)
rmse_rf = np.sqrt(mean_squared_error(y_test, y_pred_rf))
r2_rf = r2_score(y_test, y_pred_rf)

print(f'RandomForest MAE: {mae_rf:.1f}')
print(f'RandomForest RMSE: {rmse_rf:.1f}')
print(f'RandomForest R2: {r2_rf:.3f}')

fig, axs = plt.subplots(1, 3, figsize=(16, 4))


axs[0].scatter(y_test, y_pred_rf, alpha=0.3)
axs[0].plot([y_test.min(), y_test.max()], [y_test.min(), y_test.max()],
            color='r', ls='--')
axs[0].set_xlabel('Actual Traffic Volume')
axs[0].set_ylabel('Predicted Traffic Volume')
axs[0].set_title('Actual vs Predicted (RandomForest)')

residuals_rf = y_test - y_pred_rf
axs[1].scatter(y_pred_rf, residuals_rf, alpha=0.3)
axs[1].axhline(0, color='r', ls='--')
axs[1].set_xlabel('Predicted Traffic Volume')
axs[1].set_ylabel('Residuals')
axs[1].set_title('Residuals vs Predicted (RandomForest)')

sns.histplot(residuals_rf, bins=40, kde=True, ax=axs[2], color='orange')
axs[2].set_xlabel('Residuals')
axs[2].set_title('Distribution of Residuals (RandomForest)')
```

```
plt.tight_layout()
plt.show()
```

RandomForest MAE: 296.8  
RandomForest RMSE: 499.4  
RandomForest  $R^2$ : 0.936



## 1.7.5 RandomForest Model Evaluation – Diagnostic Plots

### 1. Actual vs Predicted:

- The predictions are tightly clustered around the diagonal, indicating highly accurate predictions across the range of traffic volumes. - No evidence of the “band”/funnel pattern seen in the linear model — sub-populations, rush-hours, and off-peak periods are well captured.

### 2. Residuals vs Predicted:

- Residuals are randomly scattered around zero, with no visible structure or bifurcation. - The absence of systematic errors and reduction of outlier bias confirm the ability of RandomForest to model the complex non-linear patterns in urban traffic.

### 3. Distribution of Residuals:

- The residual distribution is sharply peaked around zero, close to normal. - The tails are considerably shorter and more symmetric compared to the linear model. - No more pronounced bimodality.

#### Takeaway:

The RandomForestRegressor has dramatically improved both prediction quality and the statistical behavior of errors, demonstrating the critical value of non-linear models with time-varying, multi-modal real-world data.